

Final exam

Chaeun Shin

Installing packages

```
if(!requireNamespace("tidyverse")) install.packages("tidyverse")
```

```
## Loading required namespace: tidyverse
```

```
if(!requireNamespace("caret")) install.packages("caret")
```

```
## Loading required namespace: caret
```

```
if(!requireNamespace("neuralnet")) install.packages("neuralnet")
```

```
## Loading required namespace: neuralnet
```

```
if(!requireNamespace("keras")) install.packages("keras")
```

```
## Loading required namespace: keras
```

```
if(!requireNamespace("randomForest")) install.packages("randomForest")
```

```
## Loading required namespace: randomForest
```

```
if(!requireNamespace("rpart")) install.packages("rpart")
```

```
if(!requireNamespace("rattle")) install.packages("rattle")
```

```
## Loading required namespace: rattle
```

```
if(!requireNamespace("kernlab")) install.packages("kernlab")
```

```
## Loading required namespace: kernlab
```

Loading packages

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats   1.0.0      v stringr   1.5.1
```

```
## v ggplot2    3.5.1      v tibble    3.2.1
```

```
## v lubridate  1.9.3      v tidyr     1.3.1
```

```
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```

##
## The following object is masked from 'package:purrr':
##
## lift
library(neuralnet)

##
## Attaching package: 'neuralnet'
##
## The following object is masked from 'package:dplyr':
##
## compute
library(keras)
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
## combine
##
## The following object is masked from 'package:ggplot2':
##
## margin
library(rpart)
library(rattle)

## Loading required package: bitops
## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
##
## Attaching package: 'rattle'
##
## The following object is masked from 'package:randomForest':
##
## importance
library(kernlab)

##
## Attaching package: 'kernlab'
##
## The following object is masked from 'package:purrr':
##
## cross
##
## The following object is masked from 'package:ggplot2':
##
## alpha

```

```

library(MASS)

##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##      select
library(tidyverse)
library(glmnet)

## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following object is masked from 'package:bitops':
##
##      %&%
##
## The following objects are masked from 'package:tidyr':
##
##      expand, pack, unpack
##
## Loaded glmnet 4.1-8
library(leaps)
library(ggplot2)

```

Part I

Loading data, splitting test and training data

```

data <- read.csv('rest.csv')
data <- na.omit(data)
cat("There are", dim(data)[1], "observations left.")

```

```
## There are 4601 observations left.
```

```

set.seed(123)
training.samples <- data$y %>%
  createDataPartition(p=0.75,list=FALSE)
train_data <- data[training.samples,]
test_data <- data[-training.samples,]

```

```
data$y <- as.factor(data$y)
```

```
nrow(train_data)
```

```
## [1] 3451
```

```
nrow(test_data)
```

```
## [1] 1150
```

Logistic model, prediction

```
model <- glm(y~., train_data, family=binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
probabilities <- model %>% predict(test_data,type="response")
```

```
predicted_class_lm <- ifelse(probabilities>0.5,1,0)
```

```
confusionMatrix(factor(predicted_class_lm),factor(test_data$y))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0 658 108
```

```
##           1  36 348
```

```
##
```

```
##           Accuracy : 0.8748
```

```
##           95% CI : (0.8543, 0.8934)
```

```
## No Information Rate : 0.6035
```

```
## P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.7311
```

```
##
```

```
## Mcnemar's Test P-Value : 3.285e-09
```

```
##
```

```
##           Sensitivity : 0.9481
```

```
##           Specificity : 0.7632
```

```
##           Pos Pred Value : 0.8590
```

```
##           Neg Pred Value : 0.9062
```

```
##           Prevalence : 0.6035
```

```
##           Detection Rate : 0.5722
```

```
##           Detection Prevalence : 0.6661
```

```
##           Balanced Accuracy : 0.8556
```

```
##
```

```
##           'Positive' Class : 0
```


```
##
```

Perception model, sse, 2 neurons

```
set.seed(123)
```

```
model <- neuralnet(y~., data=train_data, hidden=2,err.fct="sse",linear.output=F)
```

```
plot(model, rep="best")
```



Error: 122.62262 Steps: 264

```
probabilities <- model %>% predict(test_data) %>% as.vector()
predicted_class_sse_2 <- ifelse(probabilities>0.5,1,0)
confusionMatrix(factor(predicted_class_sse_2),factor(test_data$y),positive='1')
```

Confusion Matrix and Statistics

##

Reference

Prediction 0 1

0 659 71

1 35 385

##

Accuracy : 0.9078

95% CI : (0.8896, 0.9239)

No Information Rate : 0.6035

P-Value [Acc > NIR] : < 2.2e-16

##

Kappa : 0.8048

##

McNemar's Test P-Value : 0.0006751

##

Sensitivity : 0.8443

Specificity : 0.9496

Pos Pred Value : 0.9167

Neg Pred Value : 0.9027

Prevalence : 0.3965


Detection Rate : 0.3348

Detection Prevalence : 0.3652

```
##      Balanced Accuracy : 0.8969
##
##      'Positive' Class : 1
##
```

Perception model, sse, 3 neurons

```
set.seed(123)
model <- neuralnet(y~., data=train_data, hidden=3,err.fct="sse",linear.output=F)
plot(model, rep="best")
```



Error: 116.160280 Steps: 20200

```
probabilities <- model %>% predict(test_data) %>% as.vector()
predicted_class_sse_3 <- ifelse(probabilities>0.5,1,0)
confusionMatrix(factor(predicted_class_sse_3),factor(test_data$y),positive='1')
```

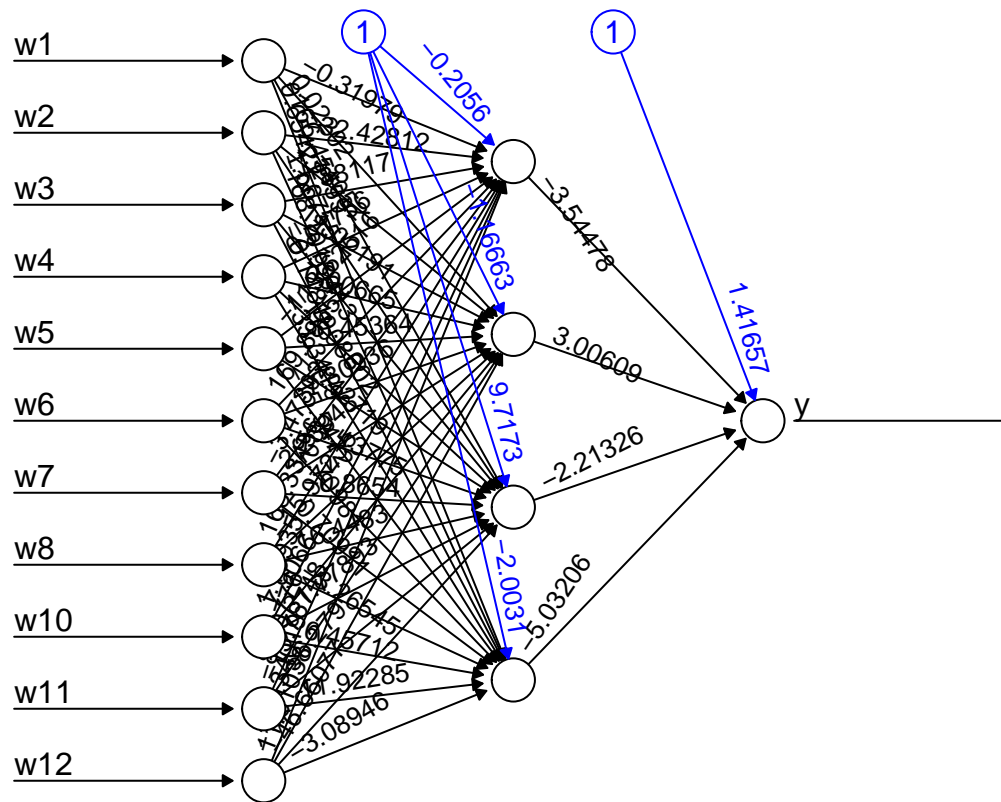
Confusion Matrix and Statistics

```
##
##      Reference
## Prediction  0   1
##      0  650  68
##      1   44 388
##
##      Accuracy : 0.9026
##      95% CI : (0.884, 0.9191)
##      No Information Rate : 0.6035
##      P-Value [Acc > NIR] : < 2e-16
##
##      Kappa : 0.7946
```

```
##
## McNemar's Test P-Value : 0.02976
##
##      Sensitivity : 0.8509
##      Specificity : 0.9366
##      Pos Pred Value : 0.8981
##      Neg Pred Value : 0.9053
##      Prevalence : 0.3965
##      Detection Rate : 0.3374
##      Detection Prevalence : 0.3757
##      Balanced Accuracy : 0.8937
##
##      'Positive' Class : 1
##
```

Perception model, sse, 4 neurons

```
set.seed(123)
model <- neuralnet(y~., data=train_data, hidden=4, err.fct="sse", linear.output=F)
plot(model, rep="best")
```



Error: 112.010812 Steps: 1000

```
probabilities <- model %>% predict(test_data) %>% as.vector()
predicted_class_sse_4 <- ifelse(probabilities>0.5,1,0)
confusionMatrix(factor(predicted_class_sse_4),factor(test_data$y),positive='1')
```

```
## Confusion Matrix and Statistics
##
##      Reference
```

```

## Prediction    0    1
##              0 655  72
##              1  39 384
##
##              Accuracy : 0.9035
##              95% CI : (0.8849, 0.9199)
##      No Information Rate : 0.6035
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.7958
##
##      McNemar's Test P-Value : 0.002387
##
##              Sensitivity : 0.8421
##              Specificity : 0.9438
##      Pos Pred Value : 0.9078
##      Neg Pred Value : 0.9010
##              Prevalence : 0.3965
##      Detection Rate : 0.3339
##      Detection Prevalence : 0.3678
##      Balanced Accuracy : 0.8930
##
##      'Positive' Class : 1
##

```

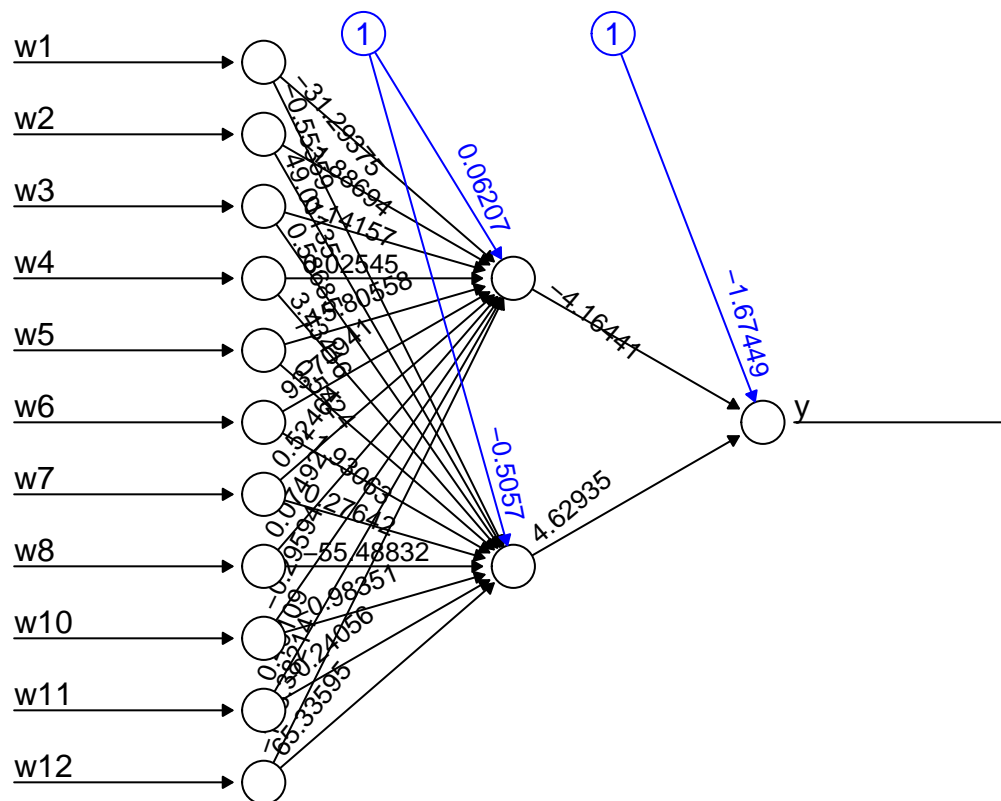
Perception model with one hidden layer with 2 neuron refers the best accuracy

Perception model, ce, 2 neurons

```

set.seed(123)
model <- neuralnet(y~., data=train_data, hidden=2,err.fct="ce",linear.output=F)
plot(model, rep="best")

```

Error: 0.76028042 Steps: 2254

```
probabilities <- model %>% predict(test_data) %>% as.vector()
predicted_class_ce_2 <- ifelse(probabilities>0.5,1,0)
confusionMatrix(factor(predicted_class_ce_2),factor(test_data$y),positive='1')
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0    1
```

```
##           0 658  73
```

```
##           1  36 383
```

```
##
```

```
##           Accuracy : 0.9052
```

```
##           95% CI : (0.8868, 0.9215)
```

```
## No Information Rate : 0.6035
```

```
## P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.7992
```

```
##
```

```
## McNemar's Test P-Value : 0.0005644
```

```
##
```

```
##           Sensitivity : 0.8399
```

```
##           Specificity : 0.9481
```

```
## Pos Pred Value : 0.9141
```

```
## Neg Pred Value : 0.9001
```

```
## Prevalence : 0.3965
```

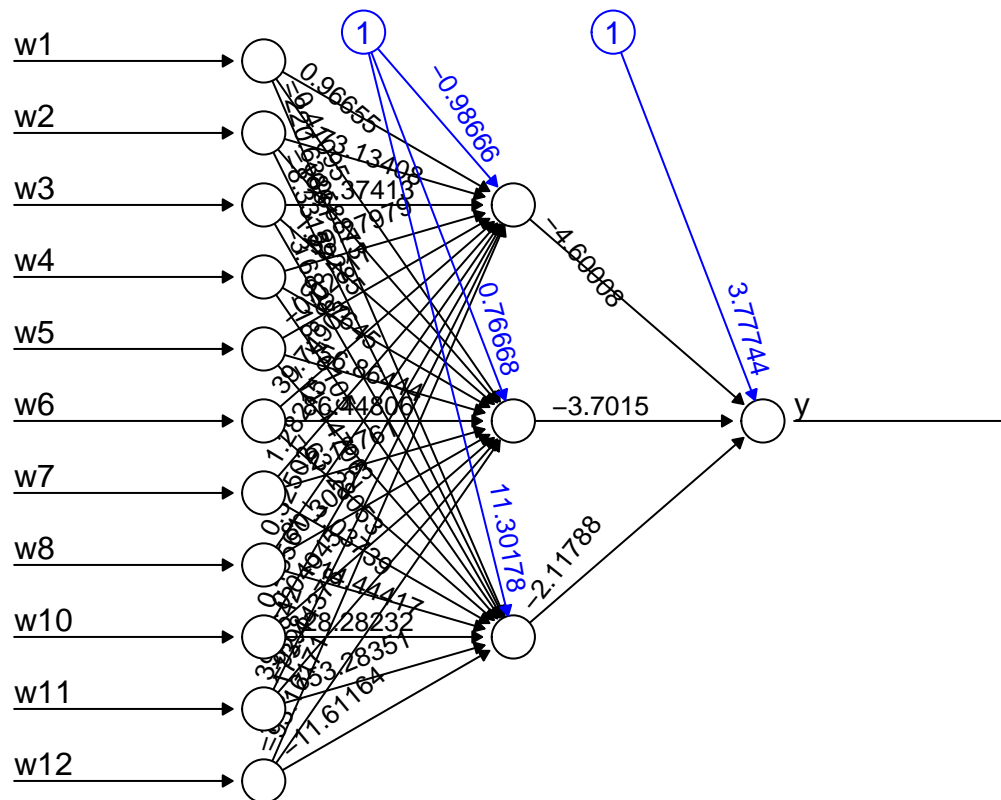
```
## Detection Rate : 0.3330
```

```
## Detection Prevalence : 0.3643
```

```
##      Balanced Accuracy : 0.8940
##
##      'Positive' Class : 1
##
```

Perception model, ce, 3 neurons

```
set.seed(123)
model <- neuralnet(y~., data=train_data, hidden=3,err.fct="ce",linear.output=F)
plot(model, rep="best")
```



Error: 822.280410 Steps: 2262

```
probabilities <- model %>% predict(test_data) %>% as.vector()
predicted_class_ce_3 <- ifelse(probabilities>0.5,1,0)
confusionMatrix(factor(predicted_class_ce_3),factor(test_data$y),positive='1')
```

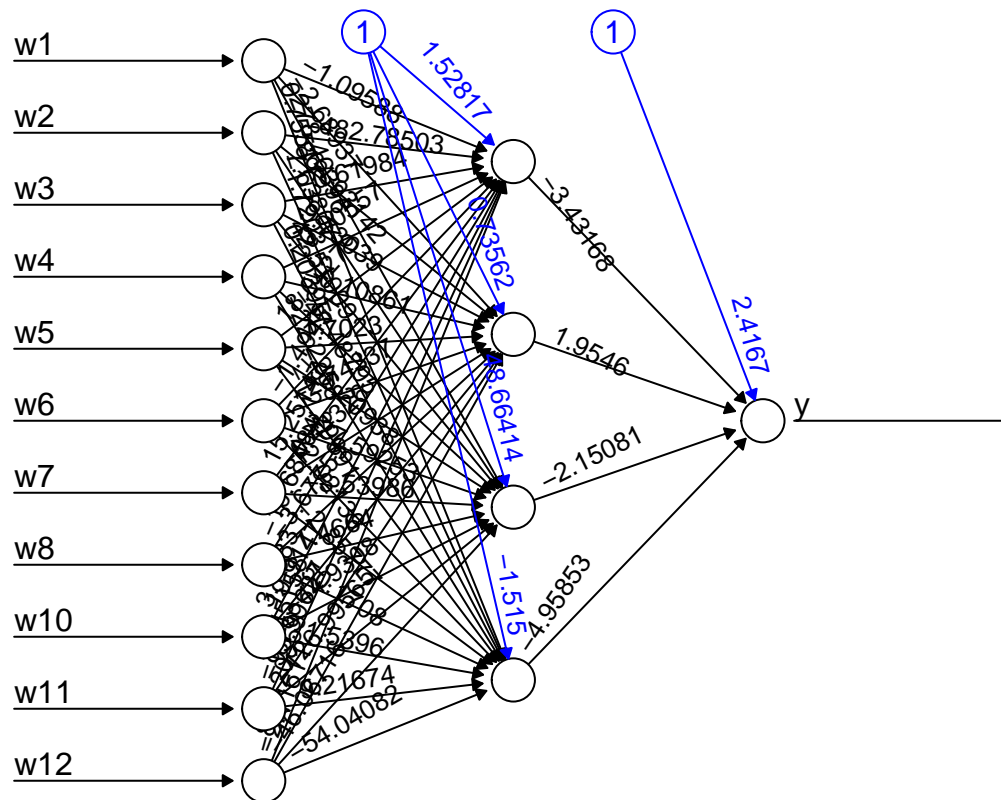
Confusion Matrix and Statistics

```
##
##      Reference
## Prediction  0   1
##      0 649  68
##      1  45 388
##
##      Accuracy : 0.9017
##      95% CI : (0.8831, 0.9183)
##      No Information Rate : 0.6035
##      P-Value [Acc > NIR] : < 2e-16
##
##      Kappa : 0.7929
```

```
##
## McNemar's Test P-Value : 0.03849
##
##      Sensitivity : 0.8509
##      Specificity : 0.9352
##      Pos Pred Value : 0.8961
##      Neg Pred Value : 0.9052
##      Prevalence : 0.3965
##      Detection Rate : 0.3374
##      Detection Prevalence : 0.3765
##      Balanced Accuracy : 0.8930
##
##      'Positive' Class : 1
##
```

Perception model, ce, 4 neurons

```
set.seed(123)
model <- neuralnet(y~., data=train_data, hidden=4, err.fct="ce", linear.output=F)
plot(model, rep="best")
```



Error: 810.560501 Steps: 9670

```
probabilities <- model %>% predict(test_data) %>% as.vector()
predicted_class_ce_4 <- ifelse(probabilities>0.5,1,0)
confusionMatrix(factor(predicted_class_ce_4),factor(test_data$y),positive='1')
```

```
## Confusion Matrix and Statistics
##
##      Reference
```

```
## Prediction    0    1
##              0 651  66
##              1  43 390
##
##              Accuracy : 0.9052
##              95% CI : (0.8868, 0.9215)
##              No Information Rate : 0.6035
##              P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.8002
##
## Mcnemar's Test P-Value : 0.0351
##
##              Sensitivity : 0.8553
##              Specificity : 0.9380
##              Pos Pred Value : 0.9007
##              Neg Pred Value : 0.9079
##              Prevalence : 0.3965
##              Detection Rate : 0.3391
##              Detection Prevalence : 0.3765
##              Balanced Accuracy : 0.8967
##
##              'Positive' Class : 1
##
```


Perception model with 2 and 4 neurons have almost same accuracy. (.9052)

SSE model provides the best accuracy for the test data.

Fully grown tree

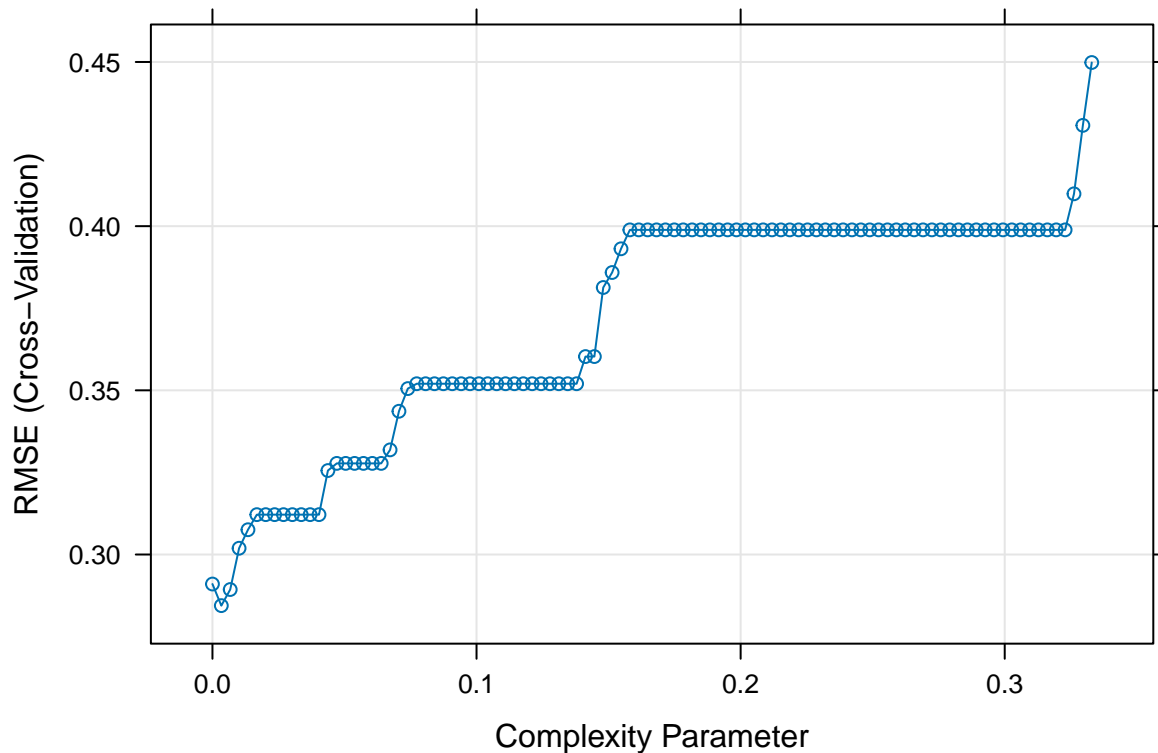
```
model <- rpart(y~., data=train_data, control=rpart.control(cp=0))
fancyRpartPlot(model)
```

Warning: labs do not fit even at cex 0.15, there may be some overplotting



Pruned tree, prediction

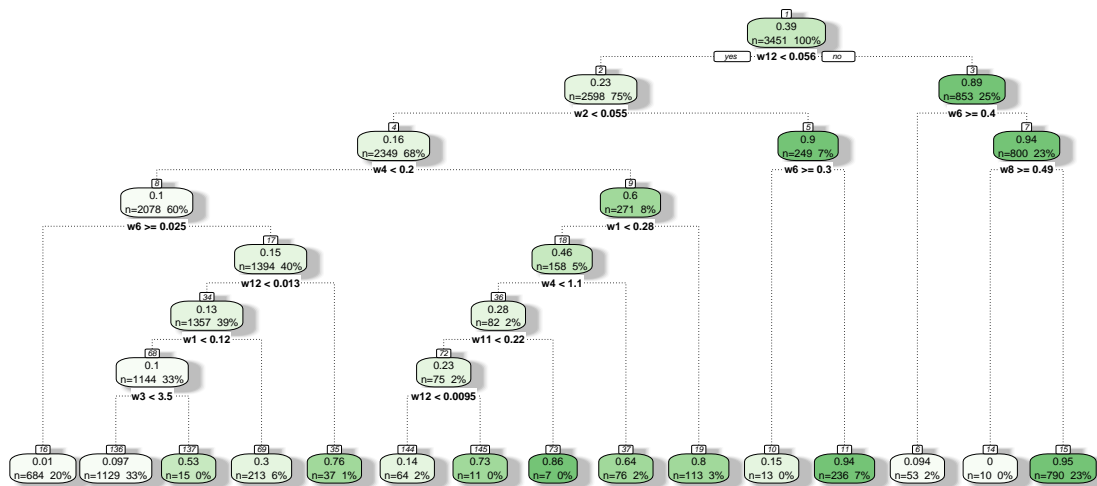
```
set.seed(123)
model2 <- train(y~., data=train_data, method="rpart", trControl=trainControl("cv", number=10), tuneLength=
## Warning in train.default(x, y, weights = w, ...): You are trying to do
## regression and your outcome only has two possible values Are you trying to do
## classification? If so, use a 2 level factor as your outcome column.
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
plot(model2)
```



```
model2$bestTune
```

```
##          cp
## 2 0.003363161
```

```
fancyRpartPlot(model2$finalModel)
```



Rattle 2024-May-14 10:15:43 chaeunshin

```
probabilities <- predict(model2, newdata=test_data)
predicted_class_prunedtree <- ifelse(probabilities>0.5,1,0)

confusionMatrix(factor(predicted_class_prunedtree),factor(test_data$y))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 645  84
##           1  49 372
##
##           Accuracy : 0.8843
##           95% CI : (0.8644, 0.9023)
##           No Information Rate : 0.6035
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.7551
##
##           McNemar's Test P-Value : 0.003197
##
##           Sensitivity : 0.9294
##           Specificity : 0.8158
##           Pos Pred Value : 0.8848
##           Neg Pred Value : 0.8836
##           Prevalence : 0.6035
##           Detection Rate : 0.5609
##           Detection Prevalence : 0.6339
##           Balanced Accuracy : 0.8726
##
##           'Positive' Class : 0
##
```

randomForest OOB

```
train_data$y <- factor(train_data$y)
test_data$y <- factor(test_data$y)
```

```
set.seed(123)
```

```
model <- train(y~., data=train_data, method="rf",trControl=trainControl("cv",number=10),importance=TRUE,
```

```
model$finalModel
```

```
##
```

```
## Call:
```

```
## randomForest(x = x, y = y, mtry = param$mtry, importance = TRUE)
```

```
##           Type of random forest: classification
```

```
##           Number of trees: 500
```

```
## No. of variables tried at each split: 6
```

```
##
```

```
##           OOB estimate of  error rate: 8.66%
```

```
## Confusion matrix:
```

```
##           0      1 class.error
```

```
## 0 1984  110  0.05253104
```

```
## 1  189 1168  0.13927782
```

Accuracy

```
(1984+1168)/(1984+110+189+1168)
```

```
## [1] 0.9133584
```

Sensitivity

```
1168/(1168+189)
```

```
## [1] 0.8607222
```

Specificity

```
1984/(1984+110)
```

```
## [1] 0.947469
```

randomForest prediction

```
predicted_class_rf <- model %>% predict(test_data)
```

```
confusionMatrix(predicted_class_rf, test_data$y, positive='1')
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0 652  74
```

```
##           1  42 382
```

```
##
```

```
##           Accuracy : 0.8991
```

```
##           95% CI : (0.8803, 0.9159)
```

```
## No Information Rate : 0.6035
```

```
## P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```


```
##           Kappa : 0.7867
```

```
##
```

```
## McNemar's Test P-Value : 0.003999
##
##      Sensitivity : 0.8377
##      Specificity : 0.9395
##      Pos Pred Value : 0.9009
##      Neg Pred Value : 0.8981
##      Prevalence : 0.3965
##      Detection Rate : 0.3322
##      Detection Prevalence : 0.3687
##      Balanced Accuracy : 0.8886
##
##      'Positive' Class : 1
##
```

Plot MeanDecreaseAccuracy


```
varImpPlot(model$finalModel, type=1)
```



Plot MeanDecreaseGini

```
varImpPlot(model$finalModel, type=2)
```


model\$finalModel




```
varImp(model,type=2)
```

```
## rf variable importance
##
##      Overall
## w12 100.000
## w2   55.443
## w4   52.016
## w6   28.509
## w1   24.673
## w10  17.670
## w3    9.704
## w8    9.668
## w5    3.101
## w11   2.502
## w7    0.000
```

Support vector machine

```
set.seed(123)
model <- train(y~., data=train_data, method="svmRadial",trControl=trainControl("cv",number=10),tuneLength=10)
plot(model)
```



```
model$bestTune
```

```
##      sigma C
## 4 0.4852082 2
```

Support vector machine prediction

```
predicted_class_svmPoly <- predict(model, newdata=test_data)
confusionMatrix(predicted_class_svmPoly, test_data$y)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##      Reference
```

```
## Prediction  0   1
```

```
##      0 647  75
```

```
##      1  47 381
```

```
##
```

```
##      Accuracy : 0.8939
```

```
##      95% CI : (0.8747, 0.9111)
```

```
##      No Information Rate : 0.6035
```

```
##      P-Value [Acc > NIR] : < 2e-16
```

```
##
```

```
##      Kappa : 0.776
```

```
##
```

```
##      McNemar's Test P-Value : 0.01451
```

```
##
```

```
##      Sensitivity : 0.9323
```

```
##      Specificity : 0.8355
```

```
##      Pos Pred Value : 0.8961
```

```
##      Neg Pred Value : 0.8902
```

```
##      Prevalence : 0.6035
```

```
##          Detection Rate : 0.5626
##    Detection Prevalence : 0.6278
##          Balanced Accuracy : 0.8839
##
##          'Positive' Class : 0
##
```

Ensemble classifier prediction

```
pred <- cbind(predicted_class_sse_2, predicted_class_rf, predicted_class_prunedtree)
pred.m <- apply(pred,1,function(x) names(which.max(table(x))))

confusionMatrix(factor(pred.m), test_data$y, positive='1')
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0    1
##          0 650  80
##          1  44 376
##
##          Accuracy : 0.8922
##          95% CI : (0.8728, 0.9095)
##    No Information Rate : 0.6035
##    P-Value [Acc > NIR] : < 2.2e-16
##
##          Kappa : 0.7716
##
##    Mcnemar's Test P-Value : 0.001672
##
##          Sensitivity : 0.8246
##          Specificity : 0.9366
##          Pos Pred Value : 0.8952
##          Neg Pred Value : 0.8904
##          Prevalence : 0.3965
##          Detection Rate : 0.3270
##    Detection Prevalence : 0.3652
##          Balanced Accuracy : 0.8806
##
##          'Positive' Class : 1
##
```

Ensemble classifier is better than the five individual classifier.


Part II

```
data <- read.csv("rest.csv")
data <- na.omit(data)
```

K-mean

```
k.means.fit <- kmeans(data,2)
library(cluster)
clusplot(data,k.means.fit$cluster,main="2D representation of the Cluster solution",color=TRUE,shade=TRUE)
```

2D representation of the Cluster solution



These two components explain 28.38 % of the point variability.

```
table(k.means.fit$cluster,data$y)
```

```
##
##          0      1
##  1 2422      7
##  2  366 1806
```

Accuracy

```
(370+7)/(366+2422+1806+7)
```

```
## [1] 0.08193871
```

Hierarchical: Ward

```
d <- dist(data,method="euclidean")
H.fit <- hclust(d,method="ward.D")
plot(H.fit)
rect.hclust(H.fit,k=2,border="red")
```

Cluster Dendrogram



d
hclust (*, "ward.D")

```
groups <- cutree(H.fit,k=2)
clusters.ward <- factor(groups,levels=1:2,labels=c(0,1))
clusplot(data,groups,main="2D representation of the Cluster solution",color=TRUE,shade=TRUE,labels=2,li
```

A PCA plot showing the relationship between Component 1 (X-axis) and Component 2 (Y-axis). The X-axis ranges from -10 to 10, and the Y-axis ranges from -10 to 20. The plot displays two main clusters of data points: a green cluster and a red cluster. The green cluster is elongated along the X-axis, while the red cluster is more compact and located at higher Component 2 values. Two shaded ellipses represent the distribution of each cluster: a blue ellipse for the green cluster and a red ellipse for the red cluster. Several points are labeled with numbers, including 1754, 2671, 2672, 2673, 2674, 2675, 2676, 2677, 2678, 2679, 2680, 2681, 2682, 2683, 2684, 2685, 2686, 2687, 2688, 2689, 2690, 2691, 2692, 2693, 2694, 2695, 2696, 2697, 2698, 2699, 2700, 2701, 2702, 2703, 2704, 2705, 2706, 2707, 2708, 2709, 2710, 2711, 2712, 2713, 2714, 2715, 2716, 2717, 2718, 2719, 2720, 2721, 2722, 2723, 2724, 2725, 2726, 2727, 2728, 2729, 2730, 2731, 2732, 2733, 2734, 2735, 2736, 2737, 2738, 2739, 2740, 2741, 2742, 2743, 2744, 2745, 2746, 2747, 2748, 2749, 2750, 2751, 2752, 2753, 2754, 2755, 2756, 2757, 2758, 2759, 2760, 2761, 2762, 2763, 2764, 2765, 2766, 2767, 2768, 2769, 2770, 2771, 2772, 2773, 2774, 2775, 2776, 2777, 2778, 2779, 2780, 2781, 2782, 2783, 2784, 2785, 2786, 2787, 2788, 2789, 2790, 2791, 2792, 2793, 2794, 2795, 2796, 2797, 2798, 2799, 2800, 2801, 2802, 2803, 2804, 2805, 2806, 2807, 2808, 2809, 2810, 2811, 2812, 2813, 2814, 2815, 2816, 2817, 2818, 2819, 2820, 2821, 2822, 2823, 2824, 2825, 2826, 2827, 2828, 2829, 2830, 2831, 2832, 2833, 2834, 2835, 2836, 2837, 2838, 2839, 2840, 2841, 2842, 2843, 2844, 2845, 2846, 2847, 2848, 2849, 2850, 2851, 2852, 2853, 2854, 2855, 2856, 2857, 2858, 2859, 2860, 2861, 2862, 2863, 2864, 2865, 2866, 2867, 2868, 2869, 2870, 2871, 2872, 2873, 2874, 2875, 2876, 2877, 2878, 2879, 2880, 2881, 2882, 2883, 2884, 2885, 2886, 2887, 2888, 2889, 2890, 2891, 2892, 2893, 2894, 2895, 2896, 2897, 2898, 2899, 2900, 2901, 2902, 2903, 2904, 2905, 2906, 2907, 2908, 2909, 2910, 2911, 2912, 2913, 2914, 2915, 2916, 2917, 2918, 2919, 2920, 2921, 2922, 2923, 2924, 2925, 2926, 2927, 2928, 2929, 2930, 2931, 2932, 2933, 2934, 2935, 2936, 2937, 2938, 2939, 2940, 2941, 2942, 2943, 2944, 2945, 2946, 2947, 2948, 2949, 2950, 2951, 2952, 2953, 2954, 2955, 2956, 2957, 2958, 2959, 2960, 2961, 2962, 2963, 2964, 2965, 2966, 2967, 2968, 2969, 2970, 2971, 2972, 2973, 2974, 2975, 2976, 2977, 2978, 2979, 2980, 2981, 2982, 2983, 2984, 2985, 2986, 2987, 2988, 2989, 2990, 2991, 2992, 2993, 2994, 2995, 2996, 2997, 2998, 2999, 3000.

```
table(data$y,clusters.ward)
```

Accuracy

```
## [1] 0.4766355
```

```
table(k.means.fit$cluster,clusters.ward)
```

Hierarchical: Average

22


Cluster Dendrogram



d
hclust (*, "average")

```
groups <- cutree(H.fit,k=2)
clusters.average <- factor(groups,levels=1:2,labels=c(0,1))
clusplot(data,groups,main="2D representation of the Cluster solution",color=TRUE,shade=TRUE,labels=2,li
```

2D representation of the Cluster solution



These two components explain 28.38 % of the point variability.

```
table(data$y,clusters.average)
```

```
##      clusters.average
##      0      1
## 0 2786      2
## 1 1812      1
```

Accuracy

```
(2786+1)/((1812+2+1+2786))
```

```
## [1] 0.6057379
```

H.Average>H.Ward>K.means based on their accuracy on confusion matrix

```
library(devtools)
```

```
## Loading required package: usethis
```

```
library(ggbiplot)
```

```
##
```

```
## Attaching package: 'ggbiplot'
```

```
## The following object is masked from 'package:rattle':
```

```
##
```

```
##      wine
```

```
data.pca <- prcomp(data,center=TRUE,scale. = TRUE)
```

```
summary(data.pca)
```


```
## Importance of components:
```

```
##      PC1      PC2      PC3      PC4      PC5      PC6      PC7
```




```
## Standard deviation      1.4173 1.1818 1.0253 1.00469 0.99147 0.97932 0.97048
## Proportion of Variance 0.1674 0.1164 0.0876 0.08412 0.08192 0.07992 0.07849
## Cumulative Proportion 0.1674 0.2838 0.3714 0.45550 0.53742 0.61734 0.69583
##              PC8      PC9      PC10      PC11      PC12
## Standard deviation    0.95448 0.91729 0.89380 0.79649 0.6814
## Proportion of Variance 0.07592 0.07012 0.06657 0.05287 0.0387
## Cumulative Proportion 0.77174 0.84186 0.90844 0.96130 1.0000
```

```
ggbiplot(data.pca)
```



```
ggbiplot(data.pca, ellipse = TRUE, groups = data$y)
```



```
data.pca$rotation[,1]
```


```
##          w1          w2          w3          w4          w5          w6          w7
## -0.2869951 -0.3584463 -0.1554061 -0.2959100 -0.1809565  0.3582133  0.2296587
##          w8          w10         w11         w12          y
##  0.1416651  0.1996604  0.0917179 -0.2843641 -0.5607341
```

The first PC contains 0.1674 of whole information of the original variable.

#Part III

```
data <- read.csv('resttt.csv')
data <- na.omit(data)
ggplot(data,aes(x=Date,y=Close)) + geom_line()
```

```
## `geom_line()`: Each group consists of only one observation.
## i Do you need to adjust the group aesthetic?
```



```
data$min_lagged <- lag(data$Low)
data$max_lagged <- log(data$High)
data$Close_norm <- (data$Close-data$min_lagged)/(data$max_lagged-data$min_lagged)

model_data <- matrix(data$Close_norm[-1])

knitr::kable(tail(model_data,10))
```

| | |
|--------|------------|
| [491,] | -0.0067433 |
| [492,] | 0.0071986 |
| [493,] | 0.0180125 |
| [494,] | -0.0343329 |
| [495,] | -0.0672942 |
| [496,] | -0.0497275 |
| [497,] | -0.0436322 |
| [498,] | -0.0254107 |
| [499,] | -0.0335002 |
| [500,] | -0.0047339 |

```
train_data <- head(model_data,-10)
test_data <- tail(model_data,10)
cat(dim(train_data)[1], 'days are divided into the training set')
```

```
## 490 days are divided into the training set
```